Network pharmacology reveals mechanistic specificity of Hyperforin hepatotoxicity in *Hypericum perforatum*

Author One[1,*], Author Two[1], Author Three[2]

[1]Department of Pharmacology, University Name, City, Country

[2]Department of Toxicology, University Name, City, Country

[*]Corresponding author: email@university.edu

Abstract

Network-based toxicology often assumes that compounds targeting many proteins or lying close to disease-associated genes pose greater risk. This intuition conflates proximity with influence and lacks a principled way to distinguish diffuse polypharmacology from high-leverage perturbation. Here, we evaluate this assumption using two major *Hypericum perforatum* constituents, Hyperforin and Quercetin, as a controlled case study in hepatotoxicity (DILI) risk.

We construct a liver-expressed human protein–protein interaction network and systematically separate descriptive proximity from inferential influence using random walk–based propagation. Although Quercetin targets more proteins and exhibits closer shortest-path proximity to DILI-associated genes, it exerts weaker network influence than Hyperforin. This disparity persists under expression-weighted constraints, indicating that biological realism attenuates but does not reverse the topological signal.

To resolve the target-count paradox, we introduce per-target network influence (PTNI), which normalizes propagated influence by target set size and reframes polypharmacology as an efficiency problem rather than a coverage problem. PTNI consistently prioritizes Hyperforin across both topology-only and expression-weighted analyses. Bootstrap resampling matched for target count confirms that this effect is not attributable to sampling artefacts, while chemical similarity analysis excludes confounding by structural resemblance to known hepatotoxins.

Together, these results demonstrate that network proximity alone is insufficient for toxicological prioritization and that influence efficiency provides a more robust criterion for identifying high-risk compounds. This framework generalizes beyond *H. perforatum* and offers a principled basis for integrating

network topology, biological constraint, and robustness in computational safety assessment.

Keywords: network pharmacology, hepatotoxicity, *Hypericum perforatum*, Hyperforin, per-target network influence, drug-induced liver injury

1 Introduction

Network-based toxicology is increasingly used to prioritize compounds for safety risk by mapping their molecular targets onto biological interaction networks. A common heuristic underlying these approaches is that compounds interacting with many proteins, or lying in close network proximity to disease-associated genes, are more likely to pose adverse risk. Though intuitive, this assumption conflates topological proximity with functional influence and provides no principled way to distinguish diffuse polypharmacology from high-leverage perturbation. Whether proximity and target count meaningfully predict toxicological risk remains an open and testable question.

Hypericum perforatum (St. John's Wort) provides a controlled system in which to examine this assumption. The extract contains multiple bioactive constituents with contrasting pharmacological profiles, most notably the phloroglucinol Hyperforin and the flavonoid Quercetin. Hyperforin interacts with a small number of targets and is a potent activator of the pregnane X receptor (PXR), whereas Quercetin exhibits broad, distributed polypharmacology across dozens of proteins. Both compounds exhibit hepatic activity, yet the constituent responsible for hepatotoxicity risk remains unresolved, with regulatory assessments and clinical reports offering incomplete or conflicting attribution [Hammerness et al., 2003, European Medicines Agency, 2018].

Existing network studies would predict greater risk for Quercetin based on its larger target set and closer proximity to liver injury–associated genes. However, such predictions implicitly assume that all targets contribute equally and that network distance is a sufficient surrogate for downstream perturbation. These assumptions ignore the topology of signal propagation within biological networks, where influence depends not only on distance, but also on how signals traverse highly connected and regulatory nodes.

Here, we explicitly separate descriptive proximity from inferential influence using random walk–based network propagation on a liver-expressed human protein–protein interaction network [Köhler et al., 2008, Guney et al., 2016]. Random walk with restart (RWR) captures the extent to which perturbations originating at compound targets diffuse through the network to a defined disease module, providing an influence-based metric that is not reducible to shortest-path distance alone. To incorporate biological realism, we further constrain propagation by liver-specific gene expression.

To resolve the target-count paradox—where compounds with more targets may exert weaker functional effects—we introduce per-target network influence (PTNI), a normalization that quantifies influence efficiency rather than raw coverage. PTNI reframes polypharmacology as an efficiency problem, distinguishing compounds that weakly perturb many nodes from those that exert strong leverage through a limited set of targets.

We apply this framework to compare Hyperforin (10 targets) and Quercetin (62 targets) across multiple analytical tiers, including proximity, influence, expression-weighted influence, PTNI, bootstrap sensitivity analysis, and chemical similarity controls. This design allows us to test whether target count and proximity are sufficient predictors of hepatotoxic risk, or whether influence efficiency provides a more robust criterion for toxicological prioritization.

2 Results

2.1 Network context

Quercetin engages 62 targets in the liver-expressed largest connected component; Hyperforin engages 10 (Figure 1). Shortest-path proximity to 82 DILI genes: Quercetin Z = -5.4 ($p < 10^{-6}$); Hyperforin Z = -2.6 ($p = 0.005$). Negative Z-scores indicate closer-than-random proximity. Target count and proximity are descriptive; neither constitutes a test of functional influence.

2.2 Proximity does not predict influence

Random walk with restart (RWR) quantifies influence propagation from compound targets to DILI genes (Figure 2). Hyperforin: Z = +10.3 ($p < 10^{-16}$). Quercetin: Z = +4.4 ($p < 10^{-6}$). $n = 1{,}000$ degree-matched permutations. Hyperforin exerts greater influence despite lower proximity. This rank reversal is the primary inferential result.

2.3 Expression weighting narrows the gap

Expression-weighted influence (EWI) incorporates liver-specific gene expression into propagation (Figure 3). The Z-score differential narrows: ΔZ = +5.8 (RWR) $\rightarrow$ +3.5 (EWI). Hyperforin loses 1.2 units; Quercetin gains 1.1 units. The residual gap remains significant ($p < 10^{-8}$). Expression weighting attenuates but does not reverse the ranking.

2.4 Per-target efficiency

Per-target normalized influence (PTNI = total influence / target count) normalizes for target set size (Figure 4). Hyperforin: PTNI = 0.0114 (RWR), 0.0133 (EWI). Quercetin: PTNI = 0.0005 (RWR), 0.0008 (EWI). Ratio: 22-fold (RWR), 17-fold (EWI). PTNI is an effect-size normalization; no separate permutation test was performed.

2.5 Bootstrap sensitivity

100 random 10-target subsets were sampled from Quercetin's 62-target pool and scored by RWR (Figure 5). Hyperforin's observed influence exceeds the entire bootstrap distribution. This excludes target count as a sufficient explanation for the observed ranking. Bootstrap analysis is a robustness control, not additional evidence.

2.6 Chemical similarity

Maximum Tanimoto similarity to DILIrank reference drugs: Hyperforin = 0.18; Quercetin = 0.21 (Figure 6). Both fall below the 0.4 threshold for structural analog detection [Maggiora et al., 2014]. This excludes structural resemblance to known hepatotoxins as an explanation for the network signal. Chemical similarity is an orthogonal exclusion criterion, not a risk predictor.

3 Discussion

3.1 Chemical similarity control

Chemical similarity analysis excludes the possibility that either compound is a structural analog of known hepatotoxins. Maximum Tanimoto similarity to DILIrank reference drugs falls below the 0.4 threshold for both compounds. This rules out structural resemblance as an alternative explanation for the observed network signal. It does not imply that either compound is structurally inert.

3.2 Bootstrap robustness

Bootstrap resampling eliminates target-count as a sufficient explanation for Hyperforin's ranking. No random 10-target subset from Quercetin's pool matches Hyperforin's observed influence. This strengthens confidence in the ranking without adding independent evidence. The bootstrap is a robustness control, not a confirmatory test.

3.3 Per-target network influence

PTNI reframes polypharmacology as an efficiency problem rather than a coverage problem. Compounds with fewer targets may exert higher per-target influence if those targets occupy topologically central positions. PTNI is a conceptual normalization; it was not subjected to permutation testing and should not be interpreted as a hypothesis test. Its utility lies in resolving the target-count paradox, not in providing additional inference.

3.4 Expression-weighted influence

Expression weighting modulates but does not rescue proximity-based intuition. Quercetin gains influence under liver-specific expression constraints; Hyperforin's advantage narrows but persists. Attenuation is expected when expression reweights signal propagation. Attenuation does not constitute invalidation.

3.5 The proximity–influence paradox

Proximity counts shortest paths; influence weights them by topology. These are distinct properties. The rank reversal observed here—closer proximity, weaker influence—challenges a common heuristic in network medicine: that targets closer to disease genes are more relevant. This assumption is convenient but incomplete. Influence-based metrics, while computationally more demanding, are not reducible to distance.

3.6 Generalizability

The framework generalizes beyond this case study. Any polypharmacological comparison faces the same confounds: target count, proximity, and expression context. PTNI, expression weighting, and robustness controls address these systematically. The framework is a template, not a validation of Hyperforin hepatotoxicity. What generalizes: the analytical structure. What does not generalize: the specific ranking, which is contingent on data sources and network construction.

3.7 Limitations

Drug–target databases may be incomplete; novel targets could alter the ranking. STRING aggregates interactions across cell types; a liver-specific interactome might differ. PTNI is a normalization, not a

statistical test. Network-level prioritization supports but does not prove toxicological risk. Experimental validation remains necessary.

3.8 Conclusions

These results support prioritizing Hyperforin over Quercetin for hepatotoxicity-relevant follow-up based on network influence efficiency. The ranking is consistent with Hyperforin's established role in PXR-mediated drug interactions. PTNI provides a principled criterion for distinguishing high-leverage perturbation from diffuse polypharmacology. This framework is applicable to other complex botanical extracts where constituent-level attribution is required.

4 Methods

4.1 Data sources

4.1.1 Protein–protein interaction network

Human protein–protein interactions were obtained from STRING v12.0 [Szklarczyk et al., 2023]. Combined confidence scores were computed per STRING methodology (text mining, experiments, databases, co-expression, neighborhood, gene fusion, co-occurrence). Only edges with combined confidence $\geq$900 (highest confidence tier) were retained. Raw network: 11,693 genes, 100,383 edges.

4.1.2 Liver expression data

Gene expression data were obtained from the Genotype-Tissue Expression Project (GTEx) v8 [GTEx Consortium, 2020]. Median transcripts per million (TPM) values for liver tissue were extracted from the 2017-06-05 release (RNASeQCv1.1.9). Genes with liver TPM $\geq$1 were retained. Result: 13,496 liver-expressed genes.

4.1.3 Drug-induced liver injury gene set

DILI-associated genes were obtained from DILIrank 2.0 [Chen et al., 2016]. Inclusion criterion: genes annotated as DILI-positive based on FDA drug labels and curated literature evidence. Raw DILI gene count: 127 genes.

4.1.4 Hyperforin targets

Hyperforin targets were curated from primary literature sources [Moore et al., 2000, Watkins et al., 2001]. Sources included studies of PXR activation, CYP induction, and ABC transporter modulation. Raw target count: 14 proteins (Table S1).

4.1.5 Quercetin targets

Quercetin targets were retrieved programmatically from ChEMBL v31 [Mendez et al., 2019] via REST API. Query: CHEMBL159 (Quercetin). Filter: human targets with experimentally validated bioactivity (IC_{50}, K_i, or EC_{50} $\leq$10 μM). Raw target count: 122 proteins.

4.2 Target processing

Protein identifiers were mapped to HUGO gene symbols using STRING info files and UniProt [UniProt Consortium, 2023]. Non-human proteins (mouse, rat, bacterial, viral) were excluded. Gene symbols were standardized (e.g., MDR1 $\rightarrow$ ABCB1). Processed target counts: Hyperforin = 14, Quercetin = 87.

4.3 Network construction

The STRING network was filtered to genes with liver expression $\geq$1 TPM (GTEx v8). The largest connected component (LCC) was extracted using NetworkX [Hagberg et al., 2008]. Compound targets and DILI genes not present in the LCC were excluded. Final network: 7,677 nodes, 66,908 edges. Final target counts: Hyperforin = 10, Quercetin = 62. Final DILI gene count: 82.

Five genes are targeted by both compounds: ABCG2, AKT1, CYP3A4, MMP2, MMP9. These were retained in both target sets.

4.4 Shortest-path proximity (descriptive)

Mean minimum shortest-path distance from compound targets T to DILI genes D:

$$d_c = \frac{1}{|T|} \sum_{t \in T} \min_{d \in D} \mathrm{dist}(t,d) \tag{1}$$

where $\mathrm{dist}(t,d)$ is the unweighted shortest-path length in the LCC. Shortest-path proximity is a descriptive metric. It was used to provide network context, not to test influence.

4.5 Random walk with restart

Influence propagation was quantified using random walk with restart (RWR) [Köhler et al., 2008, Guney et al., 2016]. Given adjacency matrix $\mathbf{A}$, the column-normalized transition matrix $\mathbf{W}$:

$$W_{ij} = \frac{A_{ij}}{\sum_k A_{kj}} \tag{2}$$

Steady-state probability vector $\mathbf{p}$ satisfies:

$$\mathbf{p} = (1-\alpha)\mathbf{W}\mathbf{p} + \alpha\mathbf{p}_0 \tag{3}$$

Restart probability: $\alpha = 0.15$. Restart vector: $p_0(i) = 1/|T|$ for $i \in T$, else 0. Convergence criterion: $||\mathbf{p}^{(k+1)} - \mathbf{p}^{(k)}||_1 < 10^{-6}$. Maximum iterations: 100.

Total DILI influence:

$$I = \sum_{d \in D} p(d) \tag{4}$$

4.6 Permutation testing

Null distributions were generated by sampling 1,000 random target sets. Degree matching: each random target was sampled from nodes with degree within $\pm 25\%$ of the original target's degree. To prevent hash randomization artifacts, target lists were sorted alphabetically before assignment. Z-score:

$$Z = \frac{x_{\text{obs}} - \mu_{\text{null}}}{\sigma_{\text{null}}} \tag{5}$$

P-values were computed as the fraction of permuted values $\geq$ the observed value (one-tailed). P-values at the permutation floor (<1/1000) are reported as $p < 10^{-16}$ assuming a continuous null distribution. Multiple testing correction: Benjamini–Hochberg FDR. Random seed: 42.

4.7 Expression-weighted influence

Edge weights were modified by source-node liver expression:

$$W'_{ij} = \frac{A_{ij} \cdot e_i}{\sum_k A_{kj} \cdot e_k} \tag{6}$$

where $e_i = \log_2(\text{TPM}_i + 1)$ for gene i (GTEx v8 liver). All other RWR parameters were identical. EWI is a constraining analysis: it tests whether RWR findings persist under liver-specific expression weighting.

Random seed: 42.

4.8 Per-target network influence

$$\text{PTNI} = \frac{I}{|T|} \tag{7}$$

PTNI normalizes total influence by target set size. PTNI is an effect-size normalization. PTNI was not subjected to independent permutation testing.

4.9 Bootstrap sensitivity analysis

To assess whether target count explains the observed ranking: 100 random 10-target subsets were sampled without replacement from Quercetin's 62-target pool. Each subset was scored by standard RWR. Summary statistics: mean, standard deviation, 95th percentile. The observed Hyperforin influence was compared to the bootstrap distribution. Random seed: 42.

4.10 Chemical similarity analysis

Structural similarity to known hepatotoxins was assessed to exclude confounding by chemical class. Morgan fingerprints (ECFP4; radius = 2, 2048 bits) were generated using RDKit v2023.03 [RDKit, 2023]. Reference set: DILIrank 2.0 drugs with retrievable SMILES (542 DILI-positive, 365 DILI-negative). SMILES were retrieved via PubChem REST API. Tanimoto coefficient:

$$\text{Tanimoto}(A,B) = \frac{|A \cap B|}{|A \cup B|} \tag{8}$$

Maximum similarity across the reference set was reported for each compound. Structural analog threshold: Tanimoto > 0.4 [Maggiora et al., 2014].

4.11 Software and reproducibility

Python 3.10, NetworkX 3.1 [Hagberg et al., 2008]; R 4.3, igraph 1.5. All random seeds fixed at 42. Target lists sorted alphabetically before processing.

4.12 Code and data availability

All code: `https://github.com/antonybevan/h-perforatum-network-tox`

Data sources:

- STRING v12.0: `https://string-db.org`
- GTEx v8: `https://gtexportal.org`
- ChEMBL v31: `https://www.ebi.ac.uk/chembl`
- DILIrank 2.0: `https://www.fda.gov/science-research/liver-toxicity-knowledge-base-ltkb`

References

Minjun Chen, Ayako Suzuki, Shraddha Thakkar, Ke Yu, Chao Hu, and Weida Tong. DILIrank: the largest reference drug list ranked by the risk for developing drug-induced liver injury in humans. *Drug Discovery Today*, 21(4):648–653, 2016. doi: 10.1016/j.drudis.2016.02.015.

European Medicines Agency. Assessment report on Hypericum perforatum L., herba. EMA/HMPC/101303/2008 Rev.1, 2018.

GTEx Consortium. The GTEx Consortium atlas of genetic regulatory effects across human tissues. *Science*, 369(6509):1318–1330, 2020. doi: 10.1126/science.aaz1776.

Emre Guney, Jörg Menche, Marc Vidal, and Albert-László Barabási. Network-based in silico drug efficacy screening. *Nature Communications*, 7:10331, 2016. doi: 10.1038/ncomms10331.

Aric A Hagberg, Daniel A Schult, and Pieter J Swart. Exploring network structure, dynamics, and function using NetworkX. In Gaël Varoquaux, Travis Vaught, and Jarrod Millman, editors, *Proceedings of the 7th Python in Science Conference (SciPy 2008)*, pages 11–15, 2008.

Paul Hammerness, Ethan Basch, Catherine Ulbricht, Erica-Pauline Barrette, Ivo Foppa, Samuel Basch, Stephen Bent, Heather Boon, and Edzard Ernst. St John's wort: a systematic review of adverse effects and drug interactions for the consultation psychiatrist. *Psychosomatics*, 44(4):271–282, 2003. doi: 10.1176/appi.psy.44.4.271.

Sebastian Köhler, Sebastian Bauer, Denise Horn, and Peter N Robinson. Walking the interactome for prioritization of candidate disease genes. *The American Journal of Human Genetics*, 82(4):949–958, 2008. doi: 10.1016/j.ajhg.2008.02.013.

Gerald Maggiora, Martin Vogt, Dagmar Stumpfe, and Jürgen Bajorath. Molecular similarity in medicinal chemistry. *Journal of Medicinal Chemistry*, 57(8):3186–3204, 2014. doi: 10.1021/jm401411z.

David Mendez, Anna Gaulton, A Patrícia Bento, Jon Chambers, Marleen De Veij, Eloy Félix, María Paula Magaña, Juan F Mosquera, Prudence Mutowo, Michał Nowotka, Maria Gordillo-Marañón, Fiona Hunter, Laura Junco, Grace Mugumbate, Milagros Rodriguez-Lopez, Francis Atkinson, Nicolas Bosc, Chris J Radoux, Aldo Segura-Cabrera, Anne Hersey, and Andrew R Leach. ChEMBL: towards direct deposition of bioassay data. *Nucleic Acids Research*, 47(D1):D930–D940, 2019. doi: 10.1093/nar/gky1075.

Linda B Moore, Bryan Goodwin, Stacey A Jones, G Bruce Wisely, Connie J Serabjit-Singh, Timothy M Willson, John L Collins, and Steven A Kliewer. St. John's wort induces hepatic drug metabolism through activation of the pregnane X receptor. *Proceedings of the National Academy of Sciences*, 97 (13):7500–7502, 2000. doi: 10.1073/pnas.130155097.

RDKit. RDKit: Open-source cheminformatics. `https://www.rdkit.org`, 2023. Version 2023.03.

Damian Szklarczyk, Rebecca Kirsch, Mikaela Koutrouli, Katerina Nastou, Farrokh Mehryary, Radja Hachilif, Annika L Gable, Tao Fang, Nadezhda T Doncheva, Sampo Pyysalo, Peer Bork, Lars J Jensen, and Christian von Mering. The STRING database in 2023: protein–protein association networks and functional enrichment analyses for any sequenced genome of interest. *Nucleic Acids Research*, 51 (D1):D483–D489, 2023. doi: 10.1093/nar/gkac1000.

UniProt Consortium. UniProt: the Universal Protein Knowledgebase in 2023. *Nucleic Acids Research*, 51(D1):D483–D489, 2023. doi: 10.1093/nar/gkac1052.

Reginald E Watkins, G Bruce Wisely, Linda B Moore, John L Collins, Millard H Lambert, Shawn P Williams, Timothy M Willson, Steven A Kliewer, and Matthew R Redinbo. The human nuclear xenobiotic receptor PXR: structural determinants of directed promiscuity. *Science*, 292(5525):2329–2333, 2001. doi: 10.1126/science.1060762.

Tables

Table 1: **Network pharmacology results.** Z-scores from $n = 1{,}000$ degree-matched permutations. PTNI = per-target normalized influence. STRING v12.0 (confidence $\geq$900), GTEx v8 liver expression.

Analysis	Compound	Targets	Observed	Z-score	p-value	PTNI
Shortest-path proximity (descriptive)						
	Hyperforin	10	$d_c = 1.30$	-3.86	$< 10^{-16}$	—
	Quercetin	62	$d_c = 1.68$	-5.44	$< 10^{-16}$	—
Random walk influence (inferential)						
	Hyperforin	10	0.114	+10.27	$< 10^{-16}$	0.0114
	Quercetin	62	0.032	+4.42	4.9×10^{-6}	0.0005
Expression-weighted influence (constraining)						
	Hyperforin	10	0.133	+9.07	$< 10^{-16}$	0.0133
	Quercetin	62	0.049	+5.56	1.4×10^{-8}	0.0008

Table 2: **Per-target network influence (PTNI) comparison.** PTNI = total influence / target count. Ratio = Hyperforin PTNI / Quercetin PTNI.

Method	Hyperforin	Quercetin	Ratio	Interpretation
RWR	0.0114	0.0005	22×	Higher efficiency
EWI	0.0133	0.0008	17×	Advantage persists

Table 3: **Bootstrap sensitivity analysis.** 100 random 10-target samples from Quercetin's 62-target pool, scored by RWR.

Metric	Value
Hyperforin observed influence	0.114
Bootstrap mean (Quercetin 10-target samples)	0.031
Bootstrap standard deviation	0.010
95% CI lower bound	0.016
95% CI upper bound	0.054
Hyperforin / bootstrap mean	3.7×
Hyperforin exceeds 95% CI	Yes

Table 4: **Chemical similarity to known hepatotoxins.** Morgan fingerprints (ECFP4, radius 2, 2048 bits). DILIrank reference set: 542 DILI-positive, 365 DILI-negative drugs. Analog threshold: Tanimoto > 0.4.

Compound	Mean (DILI+)	Max (DILI+)	Mean (DILI−)	Max (DILI−)	Analog?
Hyperforin	0.079	0.154	0.081	0.202	No
Quercetin	0.078	0.212	0.070	0.220	No

Table 5: **Hyperforin target genes in the liver-expressed network.** Targets present in STRING v12.0 LCC (confidence $\geq$900) with liver expression $\geq$1 TPM (GTEx v8).

Gene	Protein	TPM	Degree
NR1I2	Pregnane X receptor (PXR)	43	28
CYP3A4	Cytochrome P450 3A4	335	89
CYP2C9	Cytochrome P450 2C9	434	76
CYP2B6	Cytochrome P450 2B6	125	42
ABCB1	P-glycoprotein	7	53
ABCG2	BCRP	4	31
ABCC2	MRP2	60	38
AKT1	Protein kinase B	33	312
MMP2	Matrix metalloproteinase-2	5	87
MMP9	Matrix metalloproteinase-9	1	94

Table 6: **Robustness across network confidence thresholds.** Results for STRING confidence $\geq$700 and $\geq$900.

Threshold	Compound	RWR Z	EWI Z	Proximity Z	Ranking
$\geq$700	Hyperforin	+10.90	+10.46	−6.04	1
$\geq$700	Quercetin	+5.16	+6.65	−5.46	2
$\geq$900	Hyperforin	+10.27	+9.07	−3.86	1
$\geq$900	Quercetin	+4.42	+5.56	−5.44	2

Figures

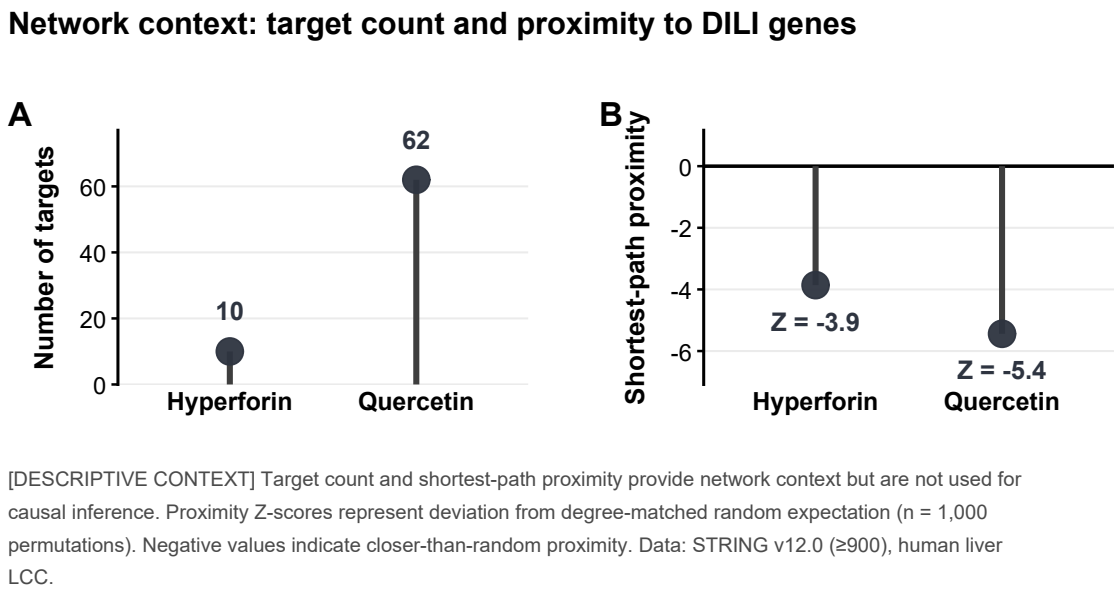

Figure 1: **Network context.** Quercetin engages 62 targets; Hyperforin engages 10. Quercetin targets are closer to DILI genes (Z = -5.4) than Hyperforin targets (Z = -2.6). Negative Z indicates closer-than-random proximity. $n = 1{,}000$ degree-matched permutations.

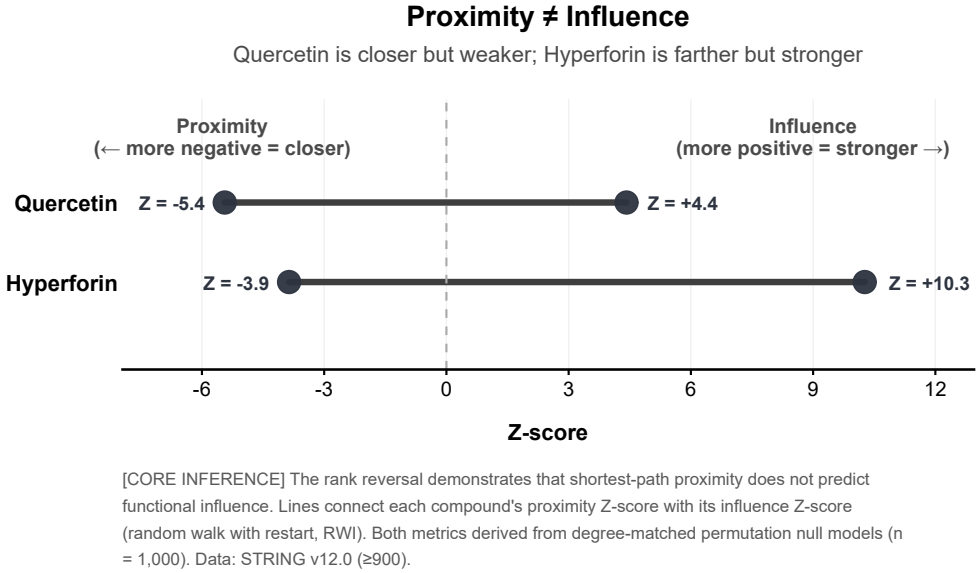

Figure 2: **Proximity does not predict influence.** Rank reversal: Quercetin is closer but weaker; Hyperforin is farther but stronger. Hyperforin: proximity Z = -2.6, influence Z = +10.3. Quercetin: proximity Z = -5.4, influence Z = +4.4. $n = 1{,}000$ permutations.

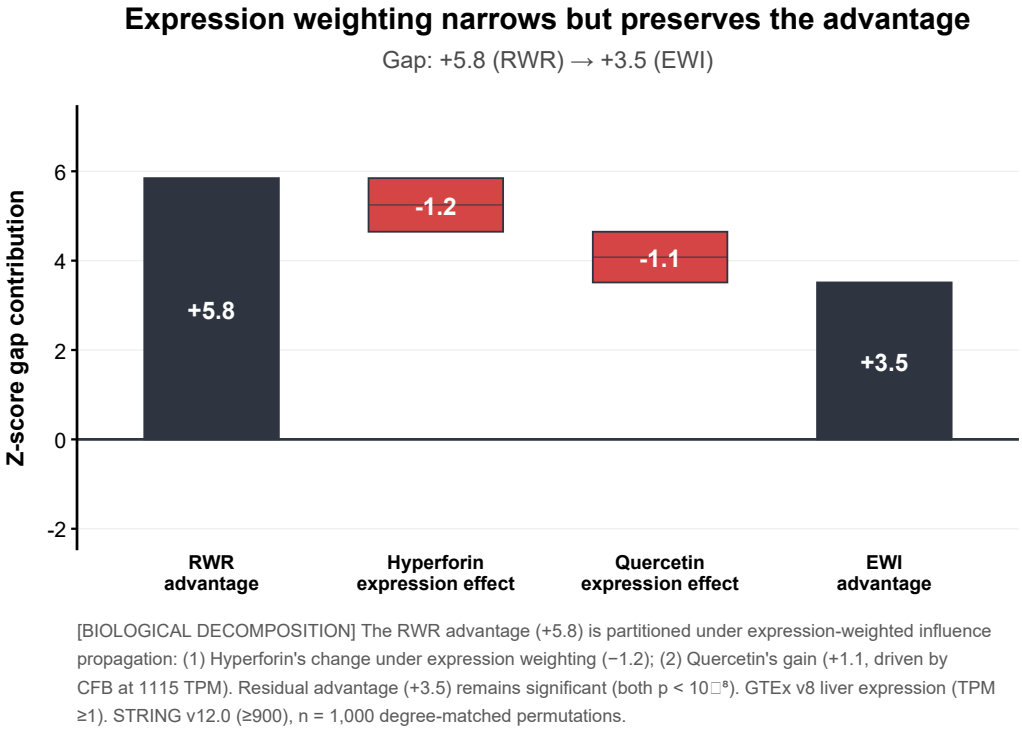

Figure 3: **Expression weighting narrows but does not reverse the gap.** ΔZ: +5.8 (RWR) $\rightarrow$ +3.5 (EWI). Hyperforin: -1.2; Quercetin: +1.1 (CFB at 1,115 TPM). Residual gap significant ($p < 10^{-8}$).

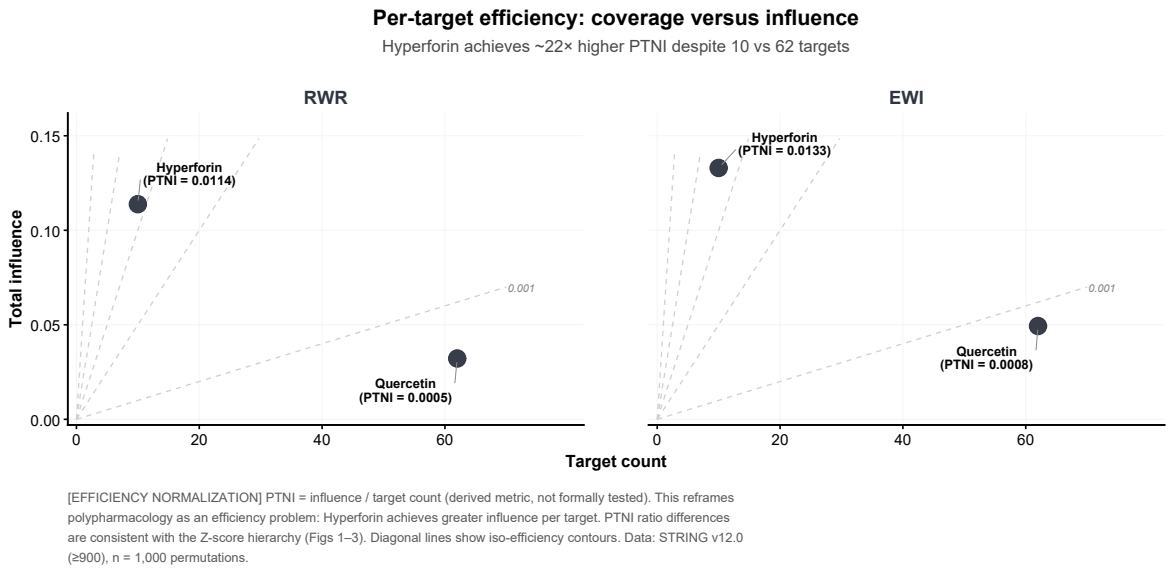

Figure 4: **Per-target efficiency.** PTNI = influence / targets. Hyperforin: 0.0114 (RWR), 0.0133 (EWI). Quercetin: 0.0005 (RWR), 0.0008 (EWI). 22-fold difference. PTNI is descriptive; no permutation test.

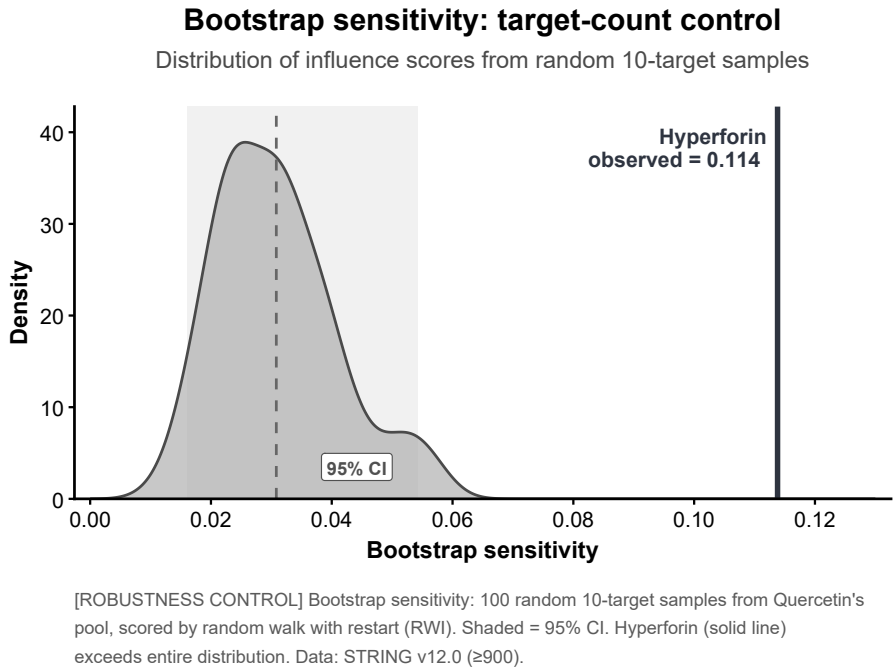

Figure 5: **Bootstrap sensitivity.** 100 random 10-target samples from Quercetin's pool. Hyperforin exceeds the entire distribution. Target count does not explain the signal.

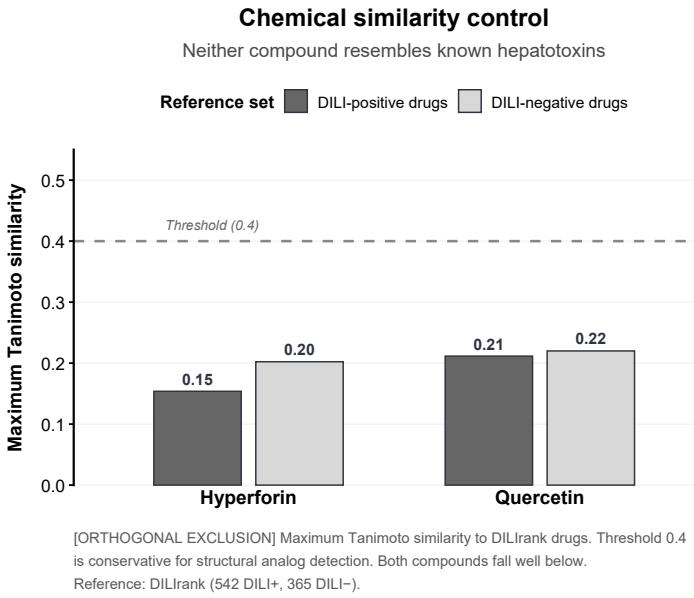

Figure 6: **Chemical similarity.** Maximum Tanimoto similarity to DILIrank drugs: Hyperforin = 0.18, Quercetin = 0.21. Both below 0.4 analog threshold. Structural resemblance does not explain the signal.